

Peer-graded Assignment: Course Final Project:

Supervised Machine Learning: Classification

About Dataset:

Context:

The dataset includes data from a random sample of 20,000 digital and 20,000 film-screen mammograms received by women age 60-89 years within the Breast Cancer Surveillance Consortium (BCSC) between January 2005 and December 2008. Some women contribute multiple examinations to the dataset. Data is useful in teaching about data analysis, epidemiological study designs, or statistical methods for binary outcomes or correlated data.

Content:

Features:

| Feature Name: | Type: | Description: |
|--|---------|--|
| Age_At_The_Time_Of_Mammography | number | Patient's age in years at time of mammogram |
| Radiologists_Assessment scale | string | Radiologist's assessment based on the BI-RADS |
| Comparison_Mammogram_From_Mammography examination available | string | Comparison mammogram from prior mammography |
| Patients_BI_RADS_Breast_Density time of mammogram | string | Patient's BI-RADS breast density as recorded at |
| Family_History_Of_Breast_Cancer relative | string | Family history of breast cancer in a first degree |
| Current_Use_Of_Hormone_Therapy mammogram | string | Current use of hormone therapy at time of |
| Binary_Indicator ever received a prior | string | Binary indicator of whether the woman had mammogram |
| History_Of_Breast_Biopsy | string | Prior history of breast biopsy |
| Is_Film_Or_Digital_Mammogra mammogram) | boolean | Film or digital mammogram (true=Digital mammogram, false=Film |

Target:

| | | |
|--|---------|---|
| Is_Binary_Indicator_Of_Cancer_Diagnosis one year of | boolean | Binary indicator of cancer diagnosis within screening mammogram (false= No cancer diagnosis, true= Cancer diagnosis) |
|--|---------|---|

Acknowledgement:

<https://www.kaggle.com/haithemhermessi/breast-cancer-screening-data-set>

Acknowledgement to Breast Cancer Surveillance Consortium (BCSC) for making this data set available for research purposes.

Objective of Analysis: Prediction of Cancer Diagnosis

```
In [1]: #Importing necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: #Importing the Data

df=pd.read_csv('data.csv',index_col='Patients_Study_ID')
```

```
In [3]: df.head()
```

```
Out[3]:
```

| | Age_At_The_Time_Of_Mammography | Radiologists_Assessment | Is_Binary_Indicator_Of_Cancer_Diagnosis | Comparison_Mammo |
|-------------------|--------------------------------|--------------------------|---|------------------|
| Patients_Study_ID | | | | |
| 1 | 62 | Negative | | False |
| 2 | 65 | Negative | | False |
| 3 | 69 | Needs additional imaging | | False |
| 4 | 64 | Benign findings | | False |
| 5 | 63 | Probably benign | | False |

```
In [4]: #List of Features

features_list=['Age_At_The_Time_Of_Mammography','Radiologists_Assessment','Is_Binary_Indicator_Of_Cancer_Diagnosis']
```

Data Cleaning and Feature Engineering:

In the given dataset, the missing values have been listed as 'Missing', hence we need to convert them to NaN values in order to impute them using central tendencies

```
In [5]: for x in features_list:
        if 'Missing' in df[x].unique():
            df[x].replace({'Missing':np.nan},inplace=True)
```

Count of missing values for each attribute is given below:

```
In [6]: df.isnull().sum()
```

```
Out[6]: Age_At_The_Time_Of_Mammography      0
Radiologists_Assessment                    0
Is_Binary_Indicator_Of_Cancer_Diagnosis    0
Comparison_Mammogram_From_Mammography    4680
Patients_BI_RADS_Breast_Density           0
Family_History_Of_Breast_Cancer           228
Current_Use_Of_Hormone_Therapy            1772
Binary_Indicator                          578
History_Of_Breast_Biopsy                  815
Is_Film_Or_Digital_Mammogram              0
Body_Mass_Index                           23208
dtype: int64
```

More than 50% values are missing for the attribute 'Body_Mass_Index', hence it can be dropped

```
In [7]: #Dropping 'Body_Mass_Index' column from dataset

df.drop(['Body_Mass_Index'],axis=1,inplace=True)
features_list.remove('Body_Mass_Index')
```

Info. about the dataset after dropping 'Body_Mass_Index':

```
In [8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 39998 entries, 1 to 36714
Data columns (total 10 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Age_At_The_Time_Of_Mammography            39998 non-null  int64
1   Radiologists_Assessment                   39998 non-null  object
2   Is_Binary_Indicator_Of_Cancer_Diagnosis   39998 non-null  bool
3   Comparison_Mammogram_From_Mammography     35318 non-null  object
4   Patients_BI_RADS_Breast_Density           39998 non-null  object
5   Family_History_Of_Breast_Cancer           39770 non-null  object
6   Current_Use_Of_Hormone_Therapy            38226 non-null  object
```

```

7   Binary_Indicator          39420 non-null object
8   History_Of_Breast_Biopsy  39183 non-null object
9   Is_Film_Or_Digital_Mammogram 39998 non-null bool
dtypes: bool(2), int64(1), object(7)
memory usage: 2.8+ MB

```

```
In [9]: sns.set(rc={'figure.figsize':(14,10)})
```

Performing Feature-Engineering:

```
In [10]: #Count-plot for 'Radiologists_Assessment' attribute

print(df['Radiologists_Assessment'].value_counts())
sns.countplot(df['Radiologists_Assessment'])


#Numerical Encoding of Categorical Values:
# 'Negative':0
# 'Probably benign':1
# 'Benign findings':2
# 'Needs additional imaging':3
# 'Suspicious abnormality':4
# 'Highly suggestive of Maglignancy':5

df['Radiologists_Assessment'].replace({'Negative':0,'Probably benign':1,'Benign findings':2,'Needs additional imaging':3,'Suspicious abnormality':4,'Highly suggestive of malignancy':5})
```

```

Negative          26031
Benign findings   10717
Needs additional imaging  3049
Probably benign    139
Suspicious abnormality  57
Highly suggestive of malignancy  5
Name: Radiologists_Assessment, dtype: int64

```



```
In [11]: #Value Counts for other attributes:

print(df['Family_History_Of_Breast_Cancer'].value_counts())

print('')

print(df['Current_Use_Of_Hormone_Therapy'].value_counts())

print('')
```

```

print(df['Binary_Indicator'].value_counts())

print('')

print(df['History_Of_Breast_Biopsy'].value_counts())

print('')

print(df['Comparison_Mammogram_From_Mammography'].value_counts())

print('')

print(df['Is_Binary_Indicator_Of_Cancer_Diagnosis'].value_counts())

print('')

print(df['Is_Film_Or_Digital_Mammogram'].value_counts())

#Binary Encoding:

df['Family_History_Of_Breast_Cancer'].replace({'Yes':1,'No':0},inplace=True)
df['Current_Use_Of_Hormone_Therapy'].replace({'Yes':1,'No':0},inplace=True)
df['Binary_Indicator'].replace({'Yes':1,'No':0},inplace=True)
df['History_Of_Breast_Biopsy'].replace({'Yes':1,'No':0},inplace=True)
df['Comparison_Mammogram_From_Mammography'].replace({'Yes':1,'No':0},inplace=True)
df['Is_Binary_Indicator_Of_Cancer_Diagnosis'].replace({'False':0,'True':1},inplace=True)
df['Is_Film_Or_Digital_Mammogram'].replace({'False':0,'True':1},inplace=True)

```

```

No      33027
Yes      6743
Name: Family_History_Of_Breast_Cancer, dtype: int64

No      33977
Yes      4249
Name: Current_Use_Of_Hormone_Therapy, dtype: int64

Yes      39124
No         296
Name: Binary_Indicator, dtype: int64

No      28733
Yes     10450
Name: History_Of_Breast_Biopsy, dtype: int64

Yes      34016
No       1302
Name: Comparison_Mammogram_From_Mammography, dtype: int64

False    39739
True       259
Name: Is_Binary_Indicator_Of_Cancer_Diagnosis, dtype: int64

True      20000
False     19998
Name: Is_Film_Or_Digital_Mammogram, dtype: int64

```

In [12]: *#Count-plot for 'Patients_BI_RADS_Breast_Density' attribute*

```

print(df['Patients_BI_RADS_Breast_Density'].value_counts())
sns.countplot(df['Patients_BI_RADS_Breast_Density'])

#Numerical Encoding of Categorical Values:
# 'Scattered fibroglandular densities':0
# 'Heterogeneously dense':1
# 'Almost entirely fatty':2
# 'Extremely dense':3

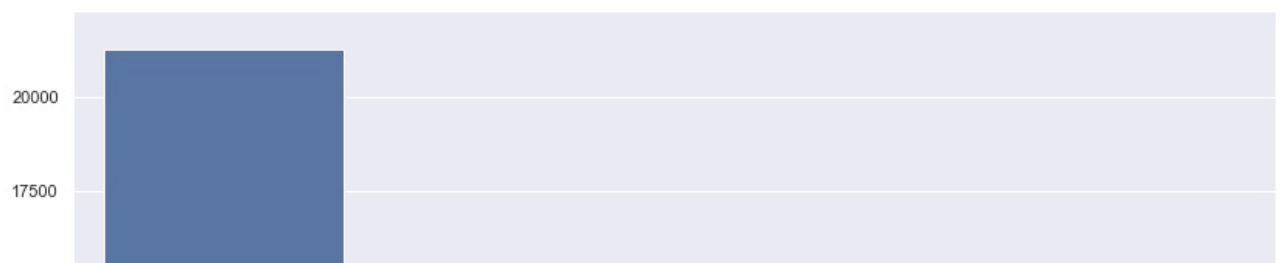
df['Patients_BI_RADS_Breast_Density'].replace({'Scattered fibroglandular densities':0,'Heterogeneously dense':1,'

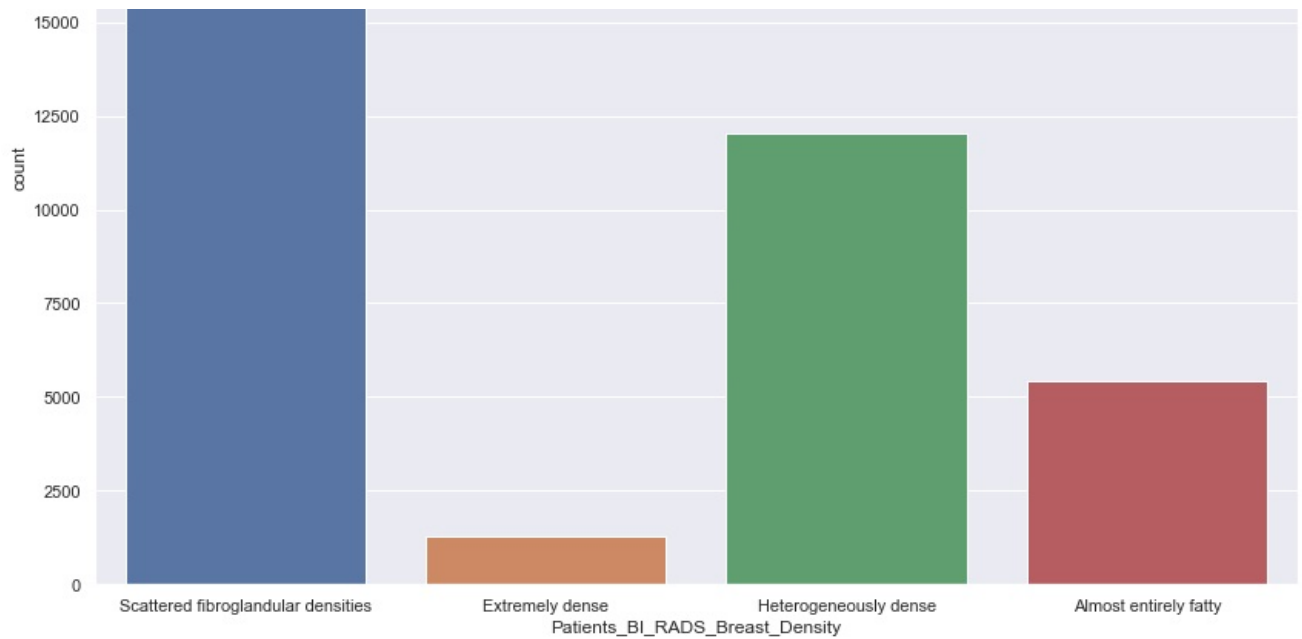
```

```

Scattered fibroglandular densities    21246
Heterogeneously dense                 12028
Almost entirely fatty                 5429
Extremely dense                       1295
Name: Patients_BI_RADS_Breast_Density, dtype: int64

```





```
In [13]: features_list=df.columns.tolist()
features_list.remove('Is_Binary_Indicator_Of_Cancer_Diagnosis')
```

Imputation of Missing Values:

```
In [14]: df.isnull().sum()
```

```
Out[14]: Age_At_The_Time_Of_Mammography      0
Radiologists_Assessment                    0
Is_Binary_Indicator_Of_Cancer_Diagnosis     0
Comparison_Mammogram_From_Mammography     4680
Patients_BI_RADS_Breast_Density            0
Family_History_Of_Breast_Cancer           228
Current_Use_Of_Hormone_Therapy            1772
Binary_Indicator                          578
History_Of_Breast_Biopsy                  815
Is_Film_Or_Digital_Mammogram              0
dtype: int64
```

```
In [15]: for x in features_list:
df[x].fillna(df[x].median(),inplace=True)
```

Missing Values for all the attributes have been imputed by the corresponding Median values

```
In [16]: features=df[features_list]
target=df['Is_Binary_Indicator_Of_Cancer_Diagnosis']
```

```
In [17]: df.isnull().sum()
```

```
Out[17]: Age_At_The_Time_Of_Mammography      0
Radiologists_Assessment                    0
Is_Binary_Indicator_Of_Cancer_Diagnosis     0
Comparison_Mammogram_From_Mammography     0
Patients_BI_RADS_Breast_Density            0
Family_History_Of_Breast_Cancer           0
Current_Use_Of_Hormone_Therapy            0
Binary_Indicator                          0
History_Of_Breast_Biopsy                  0
Is_Film_Or_Digital_Mammogram              0
dtype: int64
```

```
In [18]: df.describe()
```

```
Out[18]:
```

| | Age_At_The_Time_Of_Mammography | Radiologists_Assessment | Is_Binary_Indicator_Of_Cancer_Diagnosis | Comparison_Mammogram_From_I |
|-------|--------------------------------|-------------------------|---|-----------------------------|
| count | 39998.000000 | 39998.000000 | 39998.000000 | |
| mean | 69.555703 | 0.774364 | 0.006475 | |

| | | | |
|-----|-----------|----------|----------|
| std | 7.202581 | 1.089773 | 0.080209 |
| min | 60.000000 | 0.000000 | 0.000000 |
| 25% | 63.000000 | 0.000000 | 0.000000 |
| 50% | 68.000000 | 0.000000 | 0.000000 |
| 75% | 75.000000 | 2.000000 | 0.000000 |
| max | 89.000000 | 5.000000 | 1.000000 |

Clearly there are no outliers in the given dataset as majority of the attributes are categorical in nature

Summary of the dataset after performing Data Cleaning and Feature Engineering Techniques:

In [19]: `df.head()`

Out[19]:

| | Age_At_The_Time_Of_Mammography | Radiologists_Assessment | Is_Binary_Indicator_Of_Cancer_Diagnosis | Comparison_Mammo |
|-------------------|--------------------------------|-------------------------|---|------------------|
| Patients_Study_ID | | | | |
| 1 | 62 | 0 | 0 | |
| 2 | 65 | 0 | 0 | |
| 3 | 69 | 3 | 0 | |
| 4 | 64 | 2 | 0 | |
| 5 | 63 | 1 | 0 | |


In [20]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 39998 entries, 1 to 36714
Data columns (total 10 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Age_At_The_Time_Of_Mammography            39998 non-null  int64
1   Radiologists_Assessment                   39998 non-null  int64
2   Is_Binary_Indicator_Of_Cancer_Diagnosis   39998 non-null  int64
3   Comparison_Mammogram_From_Mammography     39998 non-null  float64
4   Patients_BI_RADS_Breast_Density           39998 non-null  int64
5   Family_History_Of_Breast_Cancer           39998 non-null  float64
6   Current_Use_Of_Hormone_Therapy            39998 non-null  float64
7   Binary_Indicator                           39998 non-null  float64
8   History_Of_Breast_Biopsy                  39998 non-null  float64
9   Is_Film_Or_Digital_Mammogram              39998 non-null  int64
dtypes: float64(5), int64(5)
memory usage: 4.6 MB
```

In [21]: `X1=features`
`y=target`

In [22]: `column_names=features_list`

In [23]: `sns.set(rc={'figure.figsize':(14,10)})`
`sns.heatmap(X1.corr(),annot=True)`





```
In [24]: sns.pairplot(X1)
```

```
Out[24]: <seaborn.axisgrid.PairGrid at 0x288e5fdb910>
```



Analysis:

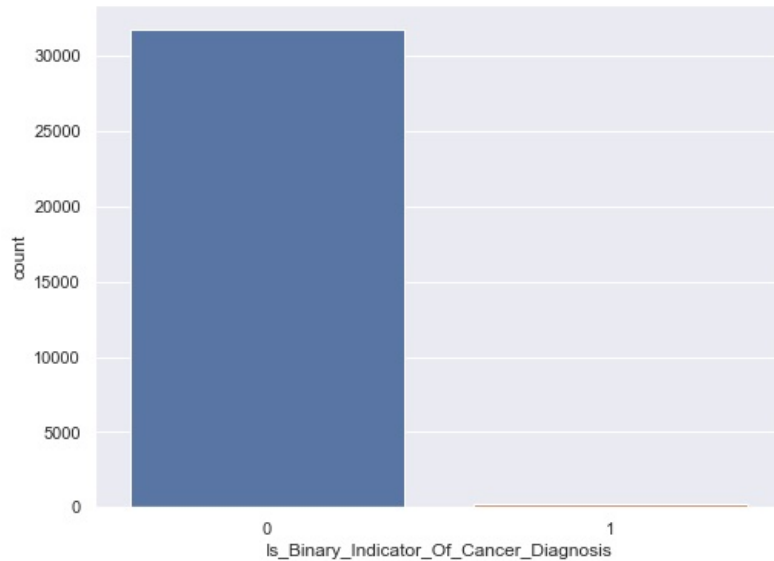
Performing Stratified Train - Test Split:

```
In [25]: from sklearn.model_selection import train_test_split
```

```
In [26]: X1_train,X_test,y1_train,y_test=train_test_split(X1,y,test_size=0.20,random_state=42,stratify=y)
```

```
In [27]: sns.set(rc={'figure.figsize':(8,6)})
sns.countplot(y1_train)
```

```
Out[27]: <AxesSubplot:xlabel='Is_Binary_Indicator_Of_Cancer_Diagnosis', ylabel='count'>
```



```
In [28]: y1_train.value_counts()
```

```
Out[28]: 0    31791
1      207
Name: Is_Binary_Indicator_Of_Cancer_Diagnosis, dtype: int64
```

It can be seen that the Target Classes are heavily imbalanced, such that No Cancer Diagnosis (0) class accounts for about >99% of the cases.

This can lead to biasing of ML models towards majority class, hence we need to either Oversample the minority class or Undersample the majority class.

Performing Oversampling of minority class using ADASYN algorithm (Adaptive Synthetic Sampling Approach):

It expands on the procedure of SMOTE, by shifting the importance of the classification boundary to those minority classes which are difficult.

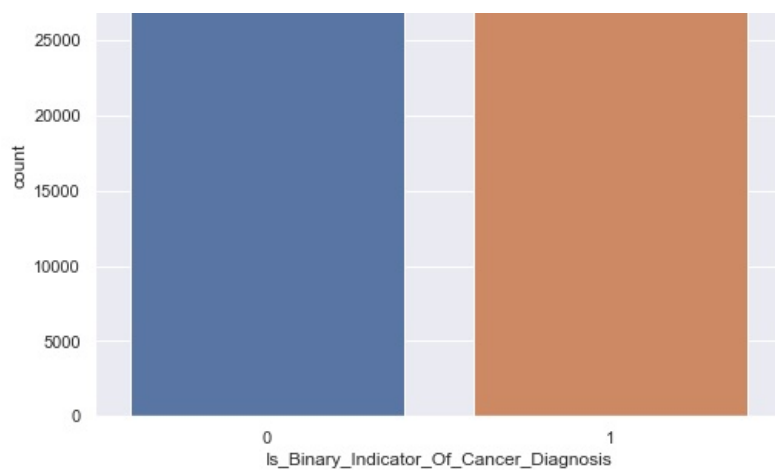
```
In [29]: from imblearn.over_sampling import ADASYN
```

```
In [30]: ada=ADASYN(random_state=42)
X2_train, y2_train = ada.fit_resample(X1_train, y1_train)
```

```
In [31]: sns.set(rc={'figure.figsize':(8,6)})
sns.countplot(y2_train)
```

```
Out[31]: <AxesSubplot:xlabel='Is_Binary_Indicator_Of_Cancer_Diagnosis', ylabel='count'>
```





```
In [32]: y2_train.value_counts()
```

```
Out[32]: 1    31840
         0    31791
         Name: Is_Binary_Indicator_Of_Cancer_Diagnosis, dtype: int64
```

Clearly both the classes have been balanced using the ADASYN technique.

However, it must be noted since we have over sampled the minority class from about ~250 cases to ~31k cases for the training data, hence we are bound to face some degree of irreducible error due to generation of such large amount of data from very small amount of data.

Note: We will be using the original data as well as oversampled data for comparison and analysis of models.

```
In [33]: sns.set(rc={'figure.figsize':(10,8)})
```

Scaling of Values for distance based algorithms:

```
In [34]: from sklearn.preprocessing import MinMaxScaler
```

```
In [35]: scaler1=MinMaxScaler()
         scaler2=MinMaxScaler()
         X1_train=scaler1.fit_transform(X1_train)
         X1_test=scaler1.transform(X_test)

         X2_train=scaler2.fit_transform(X2_train)
         X2_test=scaler2.transform(X_test)
```

Metrics to Analyse ML Models:

1. Recall Score for 'Cancer Diagnosis (1)' Class:

2. ROC Curve

We want to predict True Positive 'Cancer Diagnosis (1)' Cases, hence we want to minimise Type II Error. This can be achieved by maximising the Recall Score and AUC.

```
In [36]: from sklearn.metrics import classification_report, recall_score, roc_curve, confusion_matrix
```

```
In [37]: def report(X_train,X_test,y_train,y_test,y_train_predict,y_test_predict):
         print("Training Report:")
         rep1=classification_report(y_train,y_train_predict)
         print(rep1)
         print('Recall Score:',end=' ')
         print(recall_score(y_train,y_train_predict,average=None))
         print('')
         print('Testing Report:')
         rep2=classification_report(y_test,y_test_predict)
         print(rep2)
         print('Recall Score:',end=' ')
         print(recall_score(y_test,y_test_predict,average=None))
         print('')
         print('Confusion Matrix:')
         print(pd.DataFrame(confusion_matrix(y_test, y_test_predict)))
```

```
fpr, tpr, _ = roc_curve(y_test, y_test_predict)
#create ROC curve
plt.plot(fpr,tpr)
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```

Logistic Regression:

1. No Regularisation:

```
In [38]: from sklearn.linear_model import LogisticRegression
```

```
In [39]: lr_model=LogisticRegression(solver='liblinear',random_state=42,n_jobs=-1)
```

Original Data:

```
In [40]: lr_model.fit(X1_train,y1_train)
```

```
Out[40]: LogisticRegression(n_jobs=-1, random_state=42, solver='liblinear')
```

```
In [41]: y1_train_predict=lr_model.predict(X1_train)
y1_test_predict=lr_model.predict(X1_test)
```

```
In [42]: report(X1_train,X1_test,y1_train,y_test,y1_train_predict,y1_test_predict)
```

Training Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.99 | 1.00 | 1.00 | 31791 |
| 1 | 0.75 | 0.01 | 0.03 | 207 |
| accuracy | | | 0.99 | 31998 |
| macro avg | 0.87 | 0.51 | 0.51 | 31998 |
| weighted avg | 0.99 | 0.99 | 0.99 | 31998 |

Recall Score: [0.99996854 0.01449275]


Testing Report:


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.99 | 1.00 | 1.00 | 7948 |
| 1 | 1.00 | 0.02 | 0.04 | 52 |
| accuracy | | | 0.99 | 8000 |
| macro avg | 1.00 | 0.51 | 0.52 | 8000 |
| weighted avg | 0.99 | 0.99 | 0.99 | 8000 |

Recall Score: [1. 0.01923077]

Confusion Matrix:

| | | |
|---|------|---|
| | 0 | 1 |
| 0 | 7948 | 0 |
| 1 | 51 | 1 |





Poor and unacceptable metrics: Model is able to predict 'No Cancer Diagnosis (0)' with high precision and recall as it is the majority class and thus gets biased towards it. However we wish to achieve high recall for 'Cancer Diagnosis (1)' Class which is nearly 0 for this model and hence this model is not acceptable at all.

Oversampled Data:

```
In [43]: lr_model.fit(X2_train,y2_train)
```

```
Out[43]: LogisticRegression(n_jobs=-1, random_state=42, solver='liblinear')
```

```
In [44]: y2_train_predict=lr_model.predict(X2_train)
y2_test_predict=lr_model.predict(X2_test)
```

```
In [45]: report(X2_train,X2_test,y2_train,y_test,y2_train_predict,y2_test_predict)
```

Training Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.91 | 0.81 | 0.85 | 31791 |
| 1 | 0.83 | 0.92 | 0.87 | 31840 |
| accuracy | | | 0.86 | 63631 |
| macro avg | 0.87 | 0.86 | 0.86 | 63631 |
| weighted avg | 0.87 | 0.86 | 0.86 | 63631 |

Recall Score: [0.80758705 0.91815327]


Testing Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 0.80 | 0.89 | 7948 |
| 1 | 0.03 | 0.79 | 0.05 | 52 |
| accuracy | | | 0.80 | 8000 |
| macro avg | 0.51 | 0.80 | 0.47 | 8000 |
| weighted avg | 0.99 | 0.80 | 0.88 | 8000 |

Recall Score: [0.80259185 0.78846154]

Confusion Matrix:

| | 0 | 1 |
|---|------|------|
| 0 | 6379 | 1569 |
| 1 | 11 | 41 |





Decent metrics: ~79% Recall is decent and is a great improvement over the recall score corresponding the model for original data. Recall for the training data is ~92% which is also quite promising.

```
In [46]: from sklearn.linear_model import LogisticRegressionCV
```

2. L1 Regularisation:

```
In [47]: lr_l1=LogisticRegressionCV(Cs=30,penalty='l1',cv=6,solver='liblinear',random_state=42,scoring='recall')
```

Oversampled Data:

```
In [48]: lr_l1.fit(X2_train,y2_train)
```

```
Out[48]: LogisticRegressionCV(Cs=30, cv=6, penalty='l1', random_state=42,
                             scoring='recall', solver='liblinear')
```

```
In [49]: y2_train_predict=lr_l1.predict(X2_train)
         y2_test_predict=lr_l1.predict(X2_test)
```

```
In [50]: report(X2_train,X2_test,y2_train,y_test,y2_train_predict,y2_test_predict)
```

Training Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.90 | 0.66 | 0.76 | 31791 |
| 1 | 0.73 | 0.93 | 0.82 | 31840 |
| accuracy | | | 0.79 | 63631 |
| macro avg | 0.82 | 0.79 | 0.79 | 63631 |
| weighted avg | 0.82 | 0.79 | 0.79 | 63631 |

Recall Score: [0.65779623 0.93040201]

Testing Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 0.66 | 0.79 | 7948 |
| 1 | 0.02 | 0.87 | 0.03 | 52 |
| accuracy | | | 0.66 | 8000 |
| macro avg | 0.51 | 0.76 | 0.41 | 8000 |
| weighted avg | 0.99 | 0.66 | 0.79 | 8000 |

Recall Score: [0.65752391 0.86538462]

Confusion Matrix:

| | 0 | 1 |
|---|------|------|
| 0 | 5226 | 2722 |
| 1 | 7 | 45 |





Good and Acceptable metrics: ~87% Recall is good and is an improvement over the recall score corresponding to unregualrised model for oversampled data. Recall score for the training data is ~93% which is also quite promising.

3. L2 Regularisation:

```
In [51]: lr_l2=LogisticRegressionCV(Cs=30,penalty='l2',cv=6,solver='liblinear',random_state=42,scoring='recall')
```

Oversampled Data:

```
In [52]: lr_l2.fit(X2_train,y2_train)
```

```
Out[52]: LogisticRegressionCV(Cs=30, cv=6, random_state=42, scoring='recall',
                             solver='liblinear')
```

```
In [53]: y2_train_predict=lr_l2.predict(X2_train)
         y2_test_predict=lr_l2.predict(X2_test)
```

```
In [54]: report(X2_train,X2_test,y2_train,y_test,y2_train_predict,y2_test_predict)
```

Training Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.91 | 0.77 | 0.84 | 31791 |
| 1 | 0.80 | 0.93 | 0.86 | 31840 |
| accuracy | | | 0.85 | 63631 |
| macro avg | 0.86 | 0.85 | 0.85 | 63631 |
| weighted avg | 0.86 | 0.85 | 0.85 | 63631 |

Recall Score: [0.77477903 0.92578518]

Testing Report:


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 0.77 | 0.87 | 7948 |
| 1 | 0.02 | 0.77 | 0.04 | 52 |
| accuracy | | | 0.77 | 8000 |
| macro avg | 0.51 | 0.77 | 0.46 | 8000 |
| weighted avg | 0.99 | 0.77 | 0.86 | 8000 |

Recall Score: [0.76962758 0.76923077]

Confusion Matrix:

| | 0 | 1 |
|---|------|------|
| 0 | 6117 | 1831 |
| 1 | 12 | 40 |





Decent metrics: ~77% Recall is decent, however not an improvement over the recall score corresponding to the model for original data.

K-Nearest Neighbors:

```
In [55]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [56]: from sklearn.model_selection import GridSearchCV, StratifiedKFold
```

```
In [57]: ##Using GridSearchCV to perform CV over range of parameters and determine the best set of parameters

ss = StratifiedKFold(n_splits=6, random_state=42, shuffle=True)
n_neighbors_list=np.arange(1,50,2)
parameters = {'n_neighbors':n_neighbors_list, 'p':[1,2]}
knn=KNeighborsClassifier()
clf = GridSearchCV(knn, parameters, cv=ss, scoring='recall')
clf.fit(X2_train, y2_train)
clf.best_params_
```

```
Out[57]: {'n_neighbors': 33, 'p': 2}
```

```
In [58]: knn_model=KNeighborsClassifier(n_neighbors=clf.best_params_['n_neighbors'], p=clf.best_params_['p'])
```

Original Data:

```
In [59]: knn_model.fit(X1_train, y1_train)
```

```
Out[59]: KNeighborsClassifier(n_neighbors=33)
```

```
In [60]: y1_train_predict=knn_model.predict(X1_train)
y1_test_predict=knn_model.predict(X1_test)
```

```
In [61]: report(X1_train, X1_test, y1_train, y_test, y1_train_predict, y1_test_predict)
```

Training Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.99 | 1.00 | 1.00 | 31791 |
| 1 | 0.00 | 0.00 | 0.00 | 207 |
| accuracy | | | 0.99 | 31998 |
| macro avg | 0.50 | 0.50 | 0.50 | 31998 |
| weighted avg | 0.99 | 0.99 | 0.99 | 31998 |

Recall Score: [1. 0.]


Testing Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.99 | 1.00 | 1.00 | 7948 |
| 1 | 0.00 | 0.00 | 0.00 | 52 |
| accuracy | | | 0.99 | 8000 |
| macro avg | 0.50 | 0.50 | 0.50 | 8000 |
| weighted avg | 0.99 | 0.99 | 0.99 | 8000 |

Recall Score: [1. 0.]

Confusion Matrix:

| | | |
|---|------|---|
| | 0 | 1 |
| 0 | 7948 | 0 |
| 1 | 52 | 0 |



Poor and un-acceptable metrics: Zero recall score for 'Cancer Diagnosis (1)' Class.

Oversampled Data:

```
In [62]: knn_model.fit(X2_train,y2_train)
```

```
Out[62]: KNeighborsClassifier(n_neighbors=33)
```

```
In [63]: y2_train_predict=knn_model.predict(X2_train)
y2_test_predict=knn_model.predict(X2_test)
```

```
In [64]: report(X2_train,X2_test,y2_train,y_test,y2_train_predict,y2_test_predict)
```

Training Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.96 | 0.95 | 0.95 | 31791 |
| 1 | 0.95 | 0.96 | 0.95 | 31840 |
| accuracy | | | 0.95 | 63631 |
| macro avg | 0.95 | 0.95 | 0.95 | 63631 |
| weighted avg | 0.95 | 0.95 | 0.95 | 63631 |

Recall Score: [0.94511025 0.96092965]

Testing Report:


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 0.94 | 0.97 | 7948 |
| 1 | 0.05 | 0.52 | 0.10 | 52 |
| accuracy | | | 0.94 | 8000 |
| macro avg | 0.53 | 0.73 | 0.53 | 8000 |
| weighted avg | 0.99 | 0.94 | 0.96 | 8000 |

Recall Score: [0.94111726 0.51923077]

Confusion Matrix:

| | | |
|---|------|-----|
| | 0 | 1 |
| 0 | 7480 | 468 |
| 1 | 25 | 27 |

1.0



Unacceptable metrics: There is an improvement in Recall score as compared to model trained with original data, however it is still not acceptable.

Linear SVM:

```
In [65]: from sklearn.svm import LinearSVC
         from sklearn import svm
```

```
In [66]: LSVC=LinearSVC()
```

Original Data:

```
In [67]: LSVC.fit(X1_train,y1_train)
```

```
Out[67]: LinearSVC()
```

```
In [68]: y1_train_predict=LSVC.predict(X1_train)
         y1_test_predict=LSVC.predict(X1_test)
```

```
In [69]: report(X1_train,X1_test,y1_train,y_test,y1_train_predict,y1_test_predict)
```

Training Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.99 | 1.00 | 1.00 | 31791 |
| 1 | 0.00 | 0.00 | 0.00 | 207 |
| accuracy | | | 0.99 | 31998 |
| macro avg | 0.50 | 0.50 | 0.50 | 31998 |
| weighted avg | 0.99 | 0.99 | 0.99 | 31998 |

Recall Score: [1. 0.]

Testing Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.99 | 1.00 | 1.00 | 7948 |
| 1 | 0.00 | 0.00 | 0.00 | 52 |
| accuracy | | | 0.99 | 8000 |
| macro avg | 0.50 | 0.50 | 0.50 | 8000 |
| weighted avg | 0.99 | 0.99 | 0.99 | 8000 |


Recall Score: [1. 0.]

Confusion Matrix:

```
0 1
```



```
0 7948 0
1 52 0
```



Poor and un-acceptable metrics: Zero recall score for 'Cancer Diagnosis (1)' Class.

Oversampled Data:

```
In [70]: LSVC.fit(X2_train,y2_train)
```

```
Out[70]: LinearSVC()
```

```
In [71]: y2_train_predict=LSVC.predict(X2_train)
y2_test_predict=LSVC.predict(X2_test)
```

```
In [72]: report(X2_train,X2_test,y2_train,y_test,y2_train_predict,y2_test_predict)
```

Training Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.91 | 0.78 | 0.84 | 31791 |
| 1 | 0.81 | 0.92 | 0.86 | 31840 |
| accuracy | | | 0.85 | 63631 |
| macro avg | 0.86 | 0.85 | 0.85 | 63631 |
| weighted avg | 0.86 | 0.85 | 0.85 | 63631 |

Recall Score: [0.77751565 0.92201633]

Testing Report:


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 1.00 | 0.77 | 0.87 | 7948 |
| 1 | 0.02 | 0.81 | 0.04 | 52 |
| accuracy | | | 0.77 | 8000 |
| macro avg | 0.51 | 0.79 | 0.46 | 8000 |
| weighted avg | 0.99 | 0.77 | 0.87 | 8000 |

Recall Score: [0.7732763 0.80769231]

Confusion Matrix:

```
0 1
0 6146 1802
1 10 42
```





Decent metrics: ~81% Recall is decent and is a great improvement over the recall score corresponding the model for original data. Recall for the training data is ~92% which is also quite promising.

Stacking: Voting Classifier

```
In [73]: from sklearn.ensemble import VotingClassifier
```

Using Logistic Regression Classifier and Linear SVM via hard voting

```
In [74]: estimators=[('lr_l1',lr_l1),('linear_svm',LSVC)]
VC=VotingClassifier(estimators,voting='hard')
```

Oversampled Data:

```
In [75]: VC.fit(X2_train,y2_train)
```

```
Out[75]: VotingClassifier(estimators=[('lr_l1',
                                         LogisticRegressionCV(Cs=30, cv=6, penalty='l1',
                                                                    random_state=42,
                                                                    scoring='recall',
                                                                    solver='liblinear')),
                                      ('linear_svm', LinearSVC())])
```

```
In [76]: y2_train_predict=VC.predict(X2_train)
y2_test_predict=VC.predict(X2_test)
```

```
In [77]: report(X2_train,X2_test,y2_train,y_test,y2_train_predict,y2_test_predict)
```

Training Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.91 | 0.78 | 0.84 | 31791 |
| 1 | 0.81 | 0.92 | 0.86 | 31840 |
| accuracy | | | 0.85 | 63631 |
| macro avg | 0.86 | 0.85 | 0.85 | 63631 |
| weighted avg | 0.86 | 0.85 | 0.85 | 63631 |

Recall Score: [0.77751565 0.92201633]

Testing Report:


| | precision | recall | f1-score | support |
|---|-----------|--------|----------|---------|
| 0 | 1.00 | 0.77 | 0.87 | 7948 |
| 1 | 0.02 | 0.81 | 0.04 | 52 |

| | | | | |
|--------------|------|------|------|------|
| accuracy | | | 0.77 | 8000 |
| macro avg | 0.51 | 0.79 | 0.46 | 8000 |
| weighted avg | 0.99 | 0.77 | 0.87 | 8000 |

Recall Score: [0.7732763 0.80769231]

Confusion Matrix:

| | | |
|---|------|------|
| | 0 | 1 |
| 0 | 6146 | 1802 |
| 1 | 10 | 42 |



Decent metrics: ~81% Recall is decent. ~92% Recall for the training data is also quite promising.

Synopsis:

Objective: Prediction of Cancer Diagnosis

Here, our objective was to predict whether a person is Diagnosed with Cancer or not. We may make some mistakes in predicting a healthy person as diagnosed with cancer, however we want to minimise the error of predicting a person diagnosed with Cancer as healthy, i.e. We focused on reducing the Type II Error and on maximising the Recall Score.

Data Cleaning and Feature Engineering Techniques Used:

- 1.Imputation of Missing Data
- 2.Numerical Encoding of Categorical Data
- 3.No Outliers as almost entire dataset is categorical in nature, similarly no transformation needed
- 4.Scaling of Dataset using MinMax Scaler

Models used for Training Data:

- 1.Logistic Regression:No Regularisation (LR)
- 2.Logistic Regression:L1 Regularisation (LR_L1)
- 3.Logistic Regression:L2 Regularisation (LR_L2)
4. K-Nearest Neighbors (KNN)

5. Linear Support Vector Machine (LSVM)

6. Stacking: Voting Classifier (VC)

The Recall Scores corresponding the models are given in the below cell:

```
In [78]: print('Recall Score for Oversampled Data:')
data = {'Recall Score': ['79%', '87%', '77%', '52%', '81%', '81%']}
labels=['LR', 'LR_L1', 'LR_L2', 'KNN', 'LSVM', 'VC']
print(pd.DataFrame(data, index =labels))
```

Recall Score for Oversampled Data:

| | Recall Score |
|-------|--------------|
| LR | 79% |
| LR_L1 | 87% |
| LR_L2 | 77% |
| KNN | 52% |
| LSVM | 81% |
| VC | 81% |

Clearly Logistic Regression with L1 Regularisation is offering best Recall Score (~87%) for the target class. Hence it is most acceptable.

Note: The Classes were severely imbalanced in the original dataset, hence Oversampling techniques have been used so as to prevent biasing. Oversampling techniques do not introduce much variance, hence we are bound to face some irreducible error in our recall score.

Such models could be useful in diagnosing a patient in the absence or un-availability of a Doctor. By training upon more quality data, the model could be improved enough upon so as to be effective enough for commercial use.

PS:

The Analysis can be revisited by conducting more research and improving upon the quality of dataset.

Domain Experts can be contacted to execute the Data Cleaning and Feature Engineering tasks more effectively.

Also, the dataset could be trained upon Tree-based models to check for any improvements in Recall Score.

Analysis Conducted by-Aarohan Verma

Third Year Student at TIET, Patiala